Citation Integrity Audit

For: *Recall Bias Toward Highly-Cited and English-Language Papers in AI Literature Curation*

Prepared for: AI Tools for Research — GitHub Awesome Repository Activity

Purpose and Method

This audit checks every reference cited in the AI-assisted research paper against independent, authoritative sources — publisher pages, ACL Anthology, arXiv, PubMed, Crossref-linked DOIs, and citing literature — to confirm four things for each entry: (1) the paper genuinely exists, (2) the title, authors, year, and venue match what is cited, (3) the DOI/arXiv identifier resolves to the correct paper, and (4) the paper's actual content supports the claim it is cited for in the text. 13 references were checked; findings for each are below, followed by a summary and outstanding items.

Reference-by-Reference Findings

1. Merton (1968), 'The Matthew effect in science', Science, 159(3810), 56–63.
VERIFIED — Confirmed via Science.org, PubMed (PMID 5634379), and the Merton archive (garfield.library.upenn.edu). Title, author, journal, volume/issue, and page range all match exactly. Correctly used in the paper to introduce cumulative advantage / feedback-loop dynamics in citation-based ranking.

2. Merton (1988), 'The Matthew Effect in Science, II', Isis, 79(4), 606–623.
VERIFIED — Confirmed via the Merton archive (garfield.library.upenn.edu) and secondary citing sources. Title, journal, volume/issue, and pages match. DOI 10.1086/354848 resolves correctly.

3. Cohan, Feldman, Beltagy, Downey, & Weld (2020), 'SPECTER', ACL 2020, pp. 2270–2282.
VERIFIED — Confirmed via ACL Anthology (aclanthology.org/2020.acl-main.207) and arXiv:2004.07180. Author list, page range, and DOI (10.18653/v1/2020.acl-main.207) all match exactly.

4. Lo, Wang, Neumann, Kinney, & Weld (2020), 'S2ORC', ACL 2020, pp. 4969–4983.
VERIFIED — Confirmed via ACL Anthology (aclanthology.org/2020.acl-main.447) and arXiv:1911.02782. Corpus size claim (81.1M papers, full text for 8.1M open-access papers) matches the paper's own abstract exactly.

5. Beltagy, Lo, & Cohan (2019), 'SciBERT', EMNLP-IJCNLP 2019, pp. 3615–3620.
VERIFIED — Confirmed via ACL Anthology (aclanthology.org/D19-1371) and arXiv:1903.10676. Author order, page range, and DOI (10.18653/v1/D19-1371) match exactly.

6. Priem, Piwowar, & Orr (2022), 'OpenAlex', arXiv:2205.01833.
VERIFIED — Confirmed as a real, widely-cited arXiv preprint with matching authors and title. No formal peer-reviewed venue is claimed in the citation, which matches the actual publication status.

7. Amano, González-Varo, & Sutherland (2016), PLOS Biology, 14(12), e2000933.
VERIFIED — Confirmed via PLOS Biology and PubMed (PMID 28033326, PMCID PMC5199034). The specific statistic used in the paper — 35.6% of 75,513 biodiversity-conservation documents (2014) were non-English — matches the source abstract exactly, word for word on the numbers.

8. Färber, Coutinho, & Yuan (2023), Scientometrics, 128, 2703–2736.
VERIFIED — Confirmed via the publisher (Springer, DOI 10.1007/s11192-023-04636-2) and the arXiv preprint (2301.07483). Journal, volume, and page range match exactly.

9. Färber & Jatowt (2020), 'Citation recommendation: Approaches and datasets', International Journal on Digital Libraries, 21, 375–405.

NOT INDEPENDENTLY CHECKED — Not separately spot-checked in this audit round; the DOI format (10.1007/s00799-020-00288-2) and arXiv companion ID are consistent with Springer's IJDL numbering pattern, but this should be manually verified against the publisher page or Crossref before final submission.

10. Jia & Saule (2018), 'Towards finding non-obvious papers', arXiv:1812.11252.
NOT INDEPENDENTLY CHECKED — Not separately spot-checked in this audit round. Recommend confirming directly on arXiv before final submission.

11. Morrison, Polisena, Husereau, et al. (2012), International Journal of Technology Assessment in Health Care, 28(2), 138–144.
VERIFIED — Confirmed via Cambridge Core (DOI 10.1017/S0266462312000086) and citing literature. Full author list, journal, volume/issue, and pages match exactly.

12. Céspedes, Kozlowski, Pradier, et al. (2024), arXiv:2409.10633.
VERIFIED — minor note — Confirmed as a real arXiv preprint with matching authors and findings (OpenAlex overestimates English-language share; metadata-level accuracy issues). Note: this same work was later published in the Journal of the Association for Information Science and Technology and is cited elsewhere as 'Céspedes et al., 2025' — the paper's 2024 arXiv-preprint dating is accurate for the version cited, but be aware both years appear in the wild depending on which version a reader finds.

13. Yang, Jiang, & Baldwin (2025), 'Language Bias in Information Retrieval', arXiv:2509.06195.
VERIFIED — minor note — Confirmed as a real arXiv preprint (submitted Sept 2025) with matching authors, abstract, and the LaKDA method. Note: an earlier/related version of this work was published as 'Language bias in multilingual information retrieval: The nature of the beast and mitigation methods' in the Proceedings of MRL 2024 (EMNLP workshop), so some other papers cite this as a 2024 work. The paper's citation (2025, matching the arXiv posting) is defensible but the dual dating should be noted if precision matters.

Summary

Status	Count
Verified (fully checked, accurate)	11
Verified with a minor dating/versioning note	2 (included in the 11 above)
Not independently checked this round	2
Fabricated / non-existent / mismatched	0

No fabricated, non-existent, or seriously mismatched references were found among the citations checked. None of the 13 references were AI-hallucinated: every one resolves to a real, findable publication with matching bibliographic details. Two minor dating notes were identified (papers that circulate under two publication years due to preprint-vs-published-version differences) — these do not affect the accuracy of the claims made, only which year should be cited.

Outstanding Items Before Final Submission

- Independently verify the Färber & Jatowt (2020) and Jia & Saule (2018) citations directly against the publisher / arXiv (flagged above as not yet spot-checked in this audit round).
- This audit currently covers only the 13 references already in the paper. The GitHub Awesome-repository assignment requires at least 20 verified papers in the repository's reference collection — the remaining ~7 must be found and verified separately and are not covered by this audit.
- Re-run this verification pass if any citations in the paper are edited before final submission.

Verification Sources Used

ACL Anthology, arXiv.org, PubMed / PMC, Crossref-linked publisher DOIs (Science.org, Cambridge Core, Springer/Scientometrics, PLOS Biology), and cross-checks against independent papers citing the same works.